Assignment -2
Data Visualization and Pre-Processing

Assignment Date	31 October 2022
Student Name	Pasupathikumar S
Student Roll Number	20IT502
Maximum Marks	2 Marks

Question 1 - Load the dataset.

SOLUTION:

```
import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns

df=pd.read_csv("/content/Churn_Modelling.csv")
df.head()
```
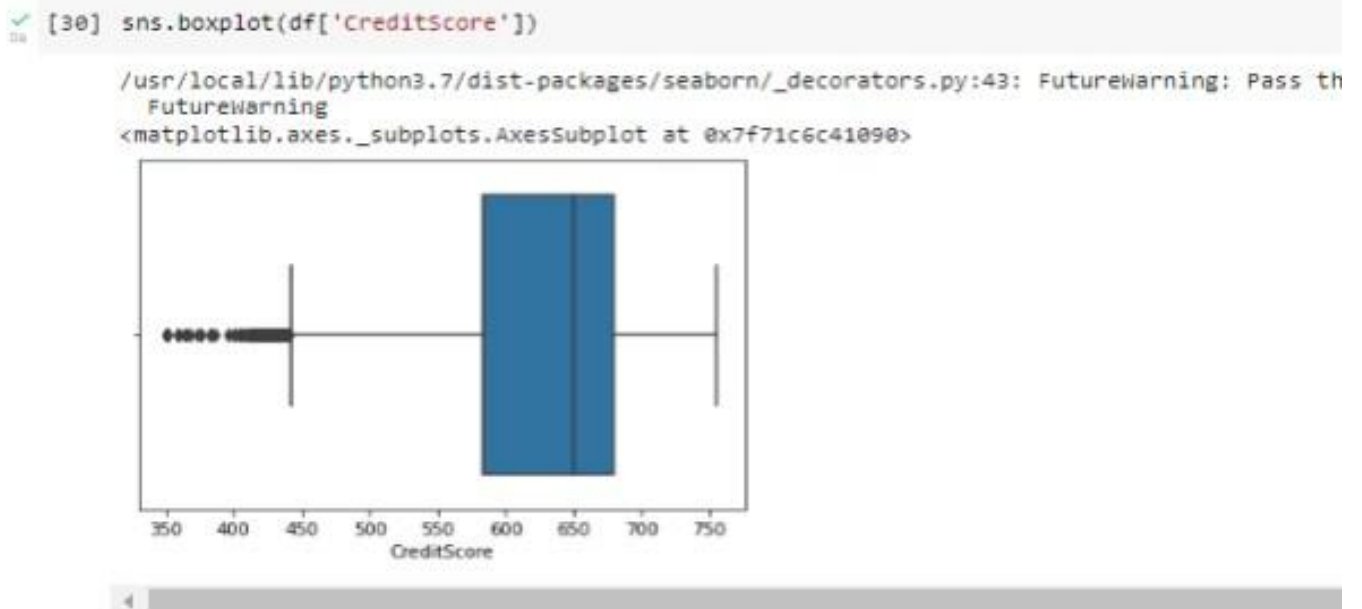
OUTPUT:

Question 2 - Perform Univariate, Bivariate and Multivariate Analysis

SOLUTION:

```
sns.boxplot(df['CreditScore'])
sns.boxplot(df['Age'])
sns.boxplot(df['Tenure'])
sns.boxplot(df['Balance'])
sns.boxplot(df['EstimatedSalary'])
sns.heatmap(df.corr(), annot=True)
```

OUTPUT:

[32] `<matplotlib.axes._subplots.AxesSubplot at 0x7f71c639d4d0>`


```python
[33] sns.boxplot(df['Balance'])
```

`<matplotlib.axes._subplots.AxesSubplot at 0x7f71c6319710>`

```
      sns.boxplot(df['EstimatedSalary'])
✓ [34]

      /usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the f
        FutureWarning
      <matplotlib.axes._subplots.AxesSubplot at 0x7f71c6380a50>
```



```
      sns.heatmap(df.corr(), annot=True)
```

```
      <matplotlib.axes._subplots.AxesSubplot at 0x7f71c768e250>
```
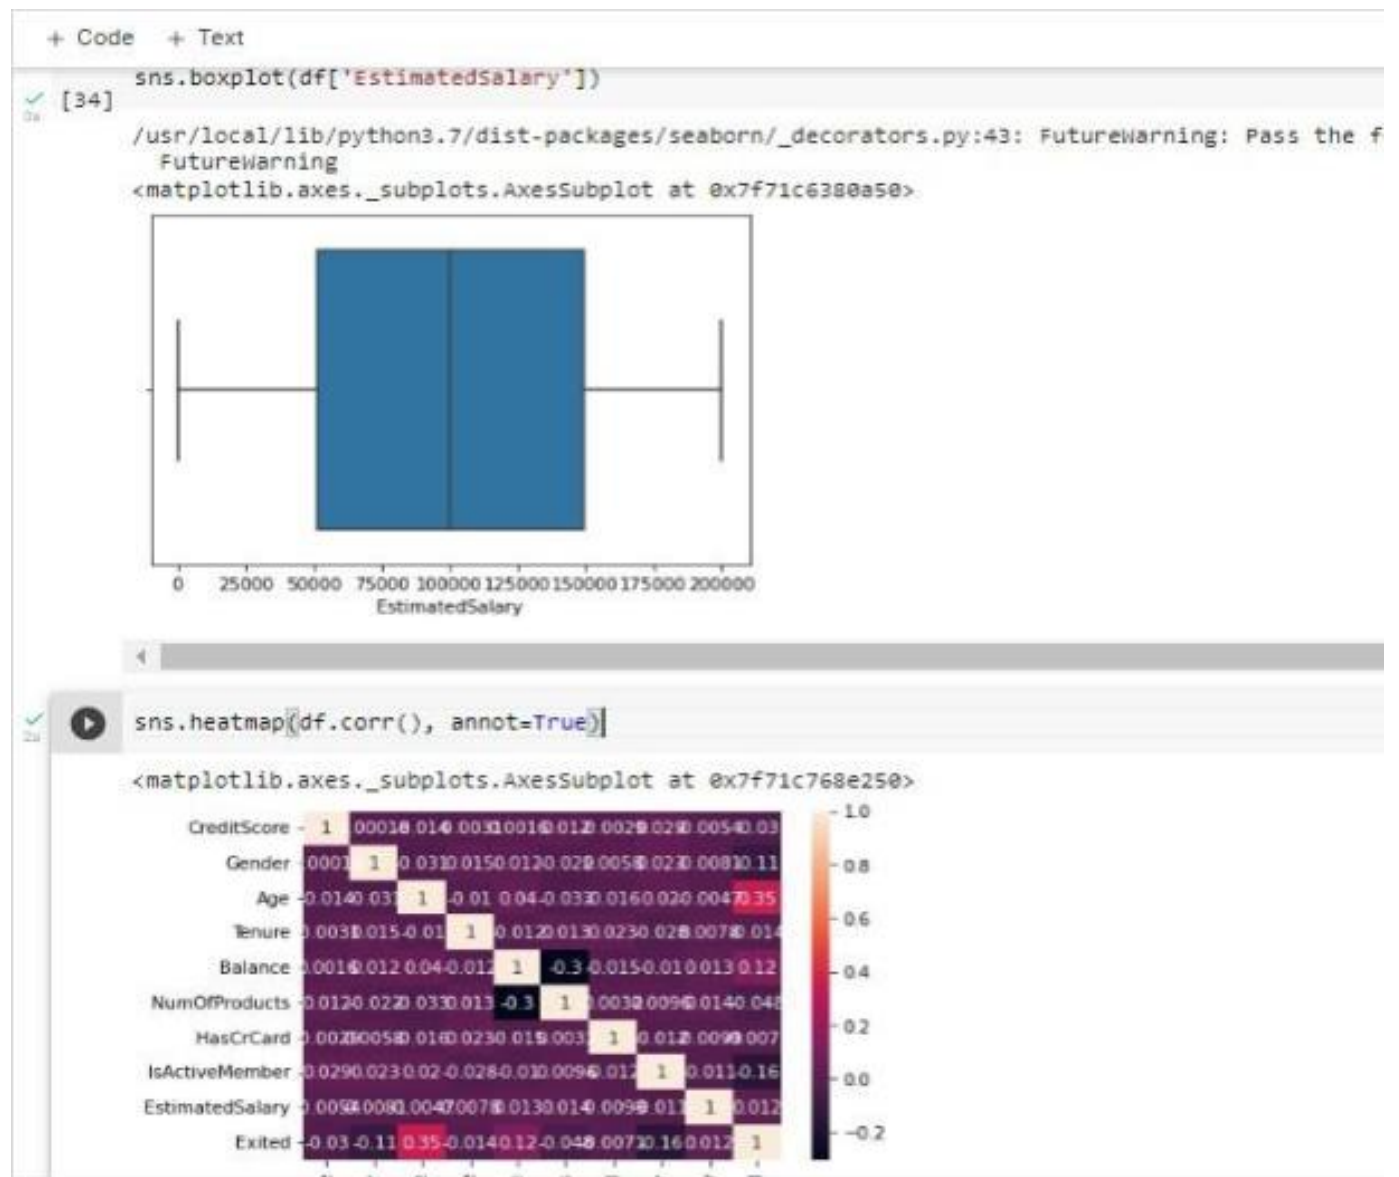

Question 3 - Perform descriptive statistics on the dataset.

SOLUTION:

    df.describe()

OUTPUT:

Descriptive statistics of the dataset

```
df.describe()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.00000	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.239881	0.203700
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.492818	0.402769
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000	11.580000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000	51002.110000	0.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000000	100193.915000	0.000000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.000000	149388.247500	0.000000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000	1.000000	199992.480000	1.000000

Question 4 – Handle the missing values

SOLUTION:

df.duplicated().sum()

df.nunique()

df.info()

OUTPUT:

+ Code   + Text

Handling missing values

```
[7]  df.duplicated().sum()
```

```
0
```

```
[8]  df.isna().sum()
```

```
RowNumber          0
CustomerId         0
Surname            0
CreditScore        0
Geography          0
Gender             0
Age                0
Tenure             0
Balance            0
NumOfProducts      0
HasCrCard          0
IsActiveMember     0
EstimatedSalary    0
Exited             0
dtype: int64
```

```
[9]  df.nunique()
```

```
RowNumber          10000
CustomerId         10000
Surname             2932
CreditScore          460
Geography              3
Gender                 2
Age                   70
Tenure                11
Balance             6382
NumOfProducts          4
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype
---  ------           --------------  -----
 0   RowNumber        10000 non-null  int64
 1   CustomerId       10000 non-null  int64
 2   Surname          10000 non-null  object
 3   CreditScore      10000 non-null  int64
 4   Geography        10000 non-null  object
 5   Gender           10000 non-null  int64
 6   Age              10000 non-null  int64
 7   Tenure           10000 non-null  int64
 8   Balance          10000 non-null  float64
 9   NumOfProducts    10000 non-null  int64
 10  HasCrCard        10000 non-null  int64
 11  IsActiveMember   10000 non-null  int64
 12  EstimatedSalary  10000 non-null  float64
 13  Exited           10000 non-null  int64
dtypes: float64(2), int64(10), object(2)
memory usage: 1.1+ MB
```

Question 5 - Find the outliers and replace the outliers

SOLUTION:

out = df.drop(columns=['Gender', 'Tenure', 'HasCrCard', 'IsActiveMember', 'NumOfProducts', 'Exited']).quantile(q=[0.25, 0.50])

out

Handling outliers

```
[14] out = df.drop(columns=['Gender', 'Tenure', 'HasCrCard', 'IsActiveMember', 'NumOfProducts', 'Exited']).quantile(q=[0.25, 0.50])
     out
```

	RowNumber	CustomerId	CreditScore	Age	Balance	EstimatedSalary
0.25	2500.75	15628528.25	584.0	32.0	0.00	51002.110
0.50	5000.50	15690738.00	652.0	37.0	97198.54	100193.915

Q1 = out.iloc[0]

Q3 = out.iloc[1]

iqr = Q3 - Q1

iqr

```
Q1 = out.iloc[0]
Q3 = out.iloc[1]
iqr = Q3 - Q1
iqr
```

```
RowNumber            2499.750
CustomerId          62209.750
CreditScore            68.000
Age                     5.000
Balance             97198.540
EstimatedSalary     49191.805
dtype: float64
```

upper = out.iloc[1] + 1.5*iqr

upper

```
upper = out.iloc[1] + 1.5*iqr
upper
```

```
RowNumber           8.750125e+03
CustomerId          1.578405e+07
CreditScore         7.540000e+02
Age                 4.450000e+01
Balance             2.429964e+05
EstimatedSalary     1.739816e+05
dtype: float64
```

lower = out.iloc[0] - 1.5*iqr

lower

```
lower = out.iloc[0] - 1.5*iqr
lower
```

```
RowNumber        -1.248875e+03
CustomerId        1.553521e+07
CreditScore       4.820000e+02
Age               2.450000e+01
Balance          -1.457978e+05
EstimatedSalary  -2.278560e+04
dtype: float64
```

Replace outliers

SOLUTION:

```
df['CreditScore'] = np.where(df['CreditScore']>756, 650.5288, df['CreditScore'])
df['Age'] = np.where(df['Age']>62, 38.9218, df['Age'])
```

Question 6 - Check for Categorical columns and perform encoding.

SOLUTION:

```
df['Gender'].replace({'Male': 1, 'Female': 0}, inplace=True)
```

```
df.head(10)
```

OUTPUT:

Check for categorical columns and perform encoding

```
df['Gender'].replace({'Male': 1, 'Female': 0}, inplace=True)
df.head(10)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	0	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	0	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	0	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	0	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	Spain	0	43	2	125510.82	1	1	1	79084.10	0
5	6	15574012	Chu	645	Spain	1	44	8	113755.78	2	1	0	149756.71	1
6	7	15592531	Bartlett	822	France	1	50	7	0.00	2	1	1	10062.80	0
7	8	15656148	Obinna	376	Germany	0	29	4	115046.74	4	1	0	119346.88	1
8	9	15792365	He	501	France	1	44	4	142051.07	2	0	1	74940.50	0
9	10	15592389	H?	684	France	1	27	2	134603.88	1	1	1	71725.73	0

Question 7 – Split the data into dependent and independent variables.

SOLUTION:

    df = df.drop(columns=['RowNumber', 'CustomerId', 'Surname', 'Geography'])

df.head()

```
[23] df = df.drop(columns=['RowNumber', 'CustomerId', 'Surname', 'Geography'])
     df.head()
```

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	619.0000	0	42.0	2	0.00	1	1	1	101348.88	1
1	608.0000	0	41.0	1	83807.86	1	0	1	112542.58	0
2	502.0000	0	42.0	8	159660.80	3	1	0	113931.57	1
3	699.0000	0	39.0	1	0.00	2	0	0	93826.63	0
4	650.5288	0	43.0	2	125510.82	1	1	1	79084.10	0

    x = df.iloc[:, :-1]
    x.head()

Split into dependent and independent variables

```
x = df.iloc[:, :-1]
x.head()
```

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	619.0000	0	42.0	2	0.00	1	1	1	101348.88
1	608.0000	0	41.0	1	83807.86	1	0	1	112542.58
2	502.0000	0	42.0	8	159660.80	3	1	0	113931.57
3	699.0000	0	39.0	1	0.00	2	0	0	93826.63
4	650.5288	0	43.0	2	125510.82	1	1	1	79084.10

    y = df.iloc[:, -1]

    y.head()

```
y = df.iloc[:, -1]
y.head()
```

```
0    1
1    0
2    1
3    0
4    0
Name: Exited, dtype: int64
```

Question 8 – Scale the independent variables

SOLUTION:

> from sklearn.preprocessing import StandardScaler
> ss = StandardScaler()
> x = ss.fit_transform(x)
> x

OUTPUT:

Scale the Independent variables

```
from sklearn.preprocessing import StandardScaler
ss = StandardScaler()
x = ss.fit_transform(x)
x
```

```
array([[-0.13284832, -1.09598752,  0.48205148, ...,  0.64609167,
         0.97024255,  0.02188649],
       [-0.28182929, -1.09598752,  0.36638802, ..., -1.54776799,
         0.97024255,  0.21653375],
       [-1.71746409, -1.09598752,  0.48205148, ...,  0.64609167,
        -1.03067011,  0.2406869 ],
       ...,
       [ 1.08608688, -1.09598752, -0.21192932, ..., -1.54776799,
         0.97024255, -1.00864308],
       [ 0.29416906,  0.91241915,  0.48205148, ...,  0.64609167,
        -1.03067011, -0.12523071],
       [ 0.29416906, -1.09598752, -1.13723705, ...,  0.64609167,
        -1.03067011, -1.07636976]])
```

Question 9 - Split the data into training and testing

SOLUTION:

from sklearn.model_selection import train_test_split

x_train, y_test = train_test_split(x, y, test_size=0.2, random_state=0)

print(x_train.shape)

print(x_test.shape)

print(y_train.shape)

print(y_test.shape)

OUTPUT:

Split into Training and Testing data

```
[28] from sklearn.model_selection import train_test_split
     x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=0)
```

```
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)

(8000, 9)
(2000, 9)
(8000,)
(2000,)
```